


Ch. 13 External Methods

External Storage


- Disk, CD-ROM, tape, ...
- 사이즈가 큰 data는 모두 internal memory에 load 될 수 없다.
 - 전체의 일부만 main memory에 있을 수 있다.
- 한 번의 disk access는 수십만 번의 instruction 수행과 맞먹는다.
- External storage access가 포함된 작업에서는 대부분 external storage access가 시간을 지배한다.
- SSD(Solid-State Drive) 시대가 열려 external storage의 부담은 많이 줄었다. But, 여전히 RAM과의 차이는 크다.



- External storage access의 단위는 block
- Typical block sizes
 - 4K, 8K, 16K, 32K bytes
 - Sector는 물리적 구조, block은 추상적 구조
 - File system 단위로 세팅/변경할 수 있다
 - Unix/Linux의 예:
 - \$ mkfs -t ext4 -b 4096 /dev/vdb1 ◀ 파일의 block size를 4K로 세팅
- Each block consists of multiple data records
- Typical access
 - `buf.readBlock(dataFile, i)` ◀ `dataFile`에서 i 번 block을 `buf`로 읽어들이м
 - `buf`: object for memory buffer
 - `dataFile`: file containing the data
 - `i`: index of block


Java

java.lang.Object
java.io.Reader
 java.io.BufferedReader
java.io.Writer
 java.io.BufferedWriter




External Mergesort

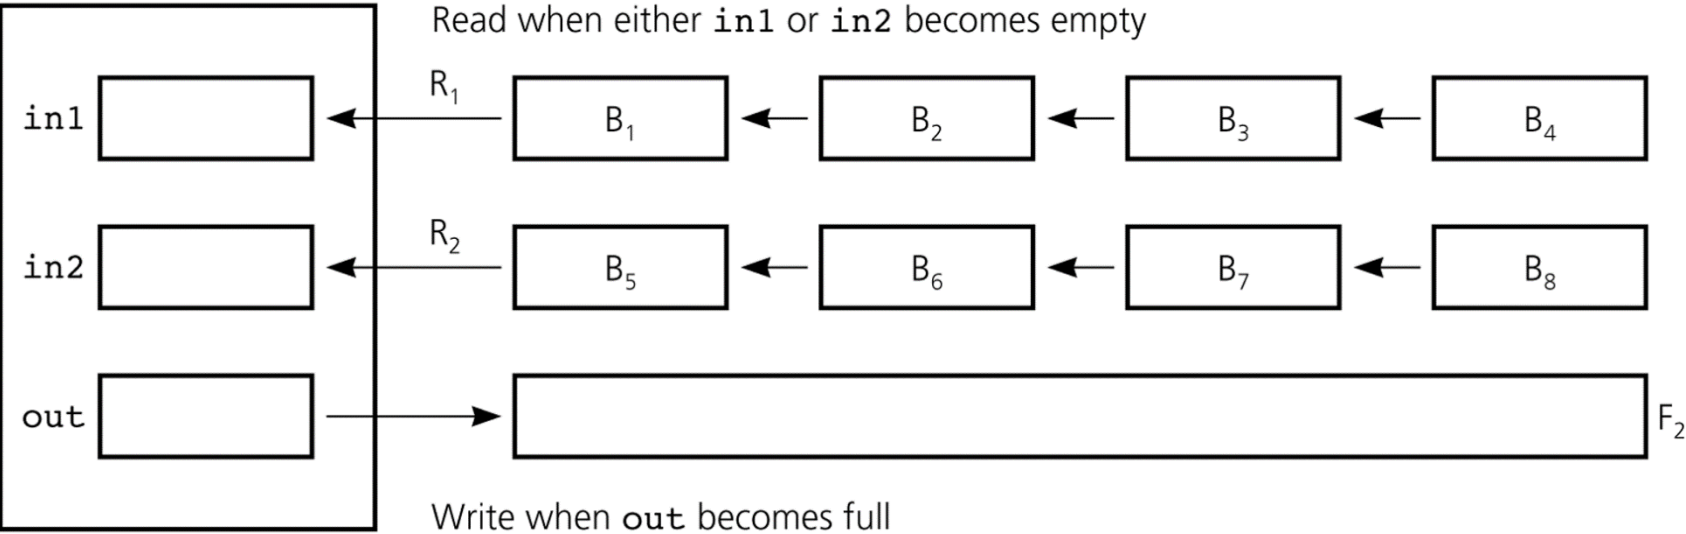
- 비현실적인 예제
 - 1,600 employee records
 - 100 records/block
 - Maximum 300 records on memory
- Files
 - F : Original data
 - F_1, F_2 : work files



Merging Single Blocks




Merging Two 4-Block Runs

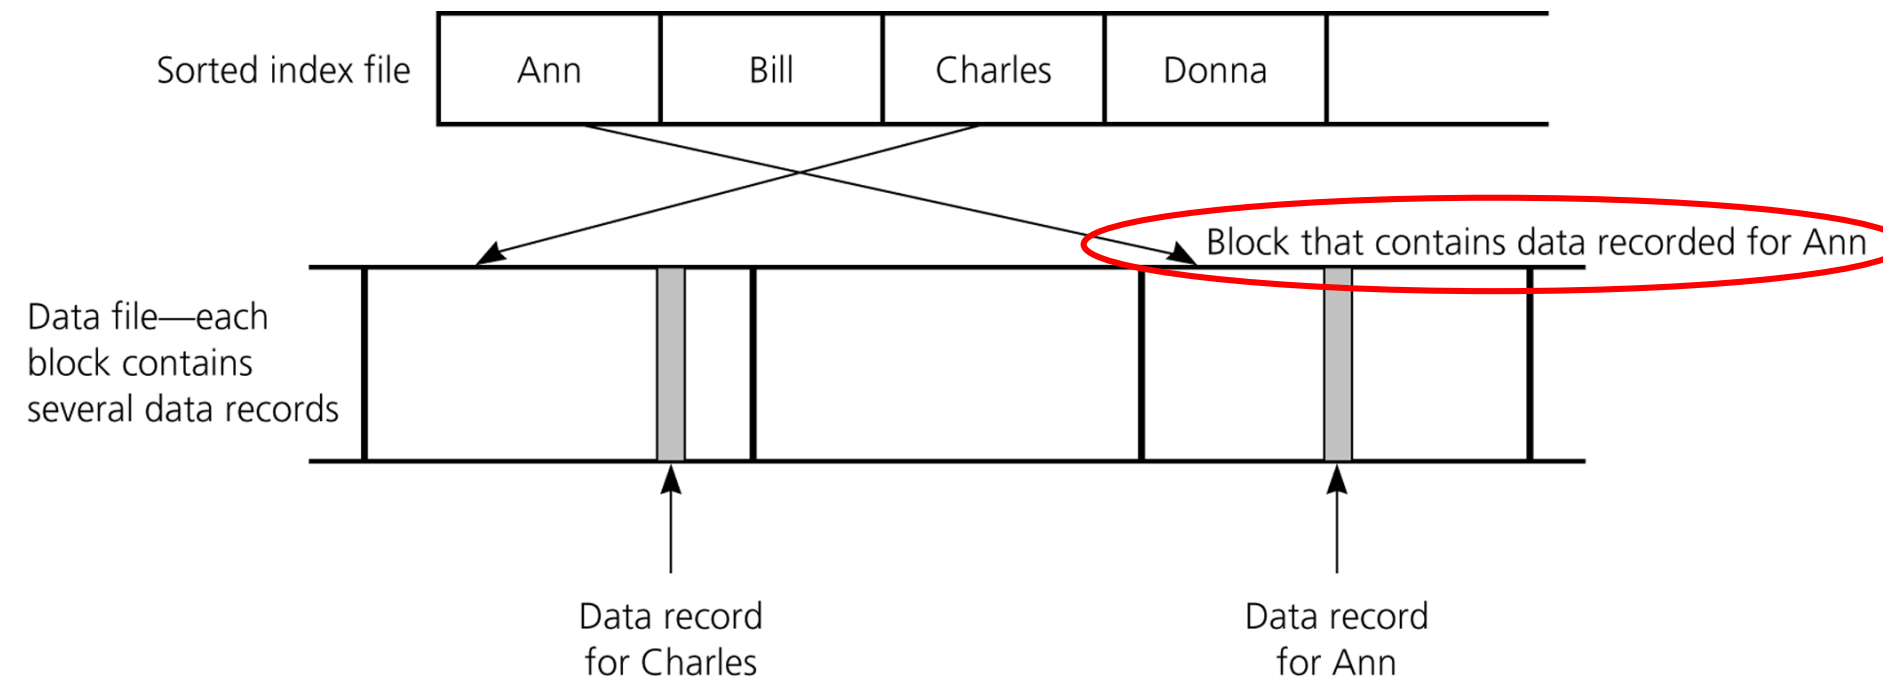


External Tables

- Data are stored in an external storage
- External hashing
- B-trees



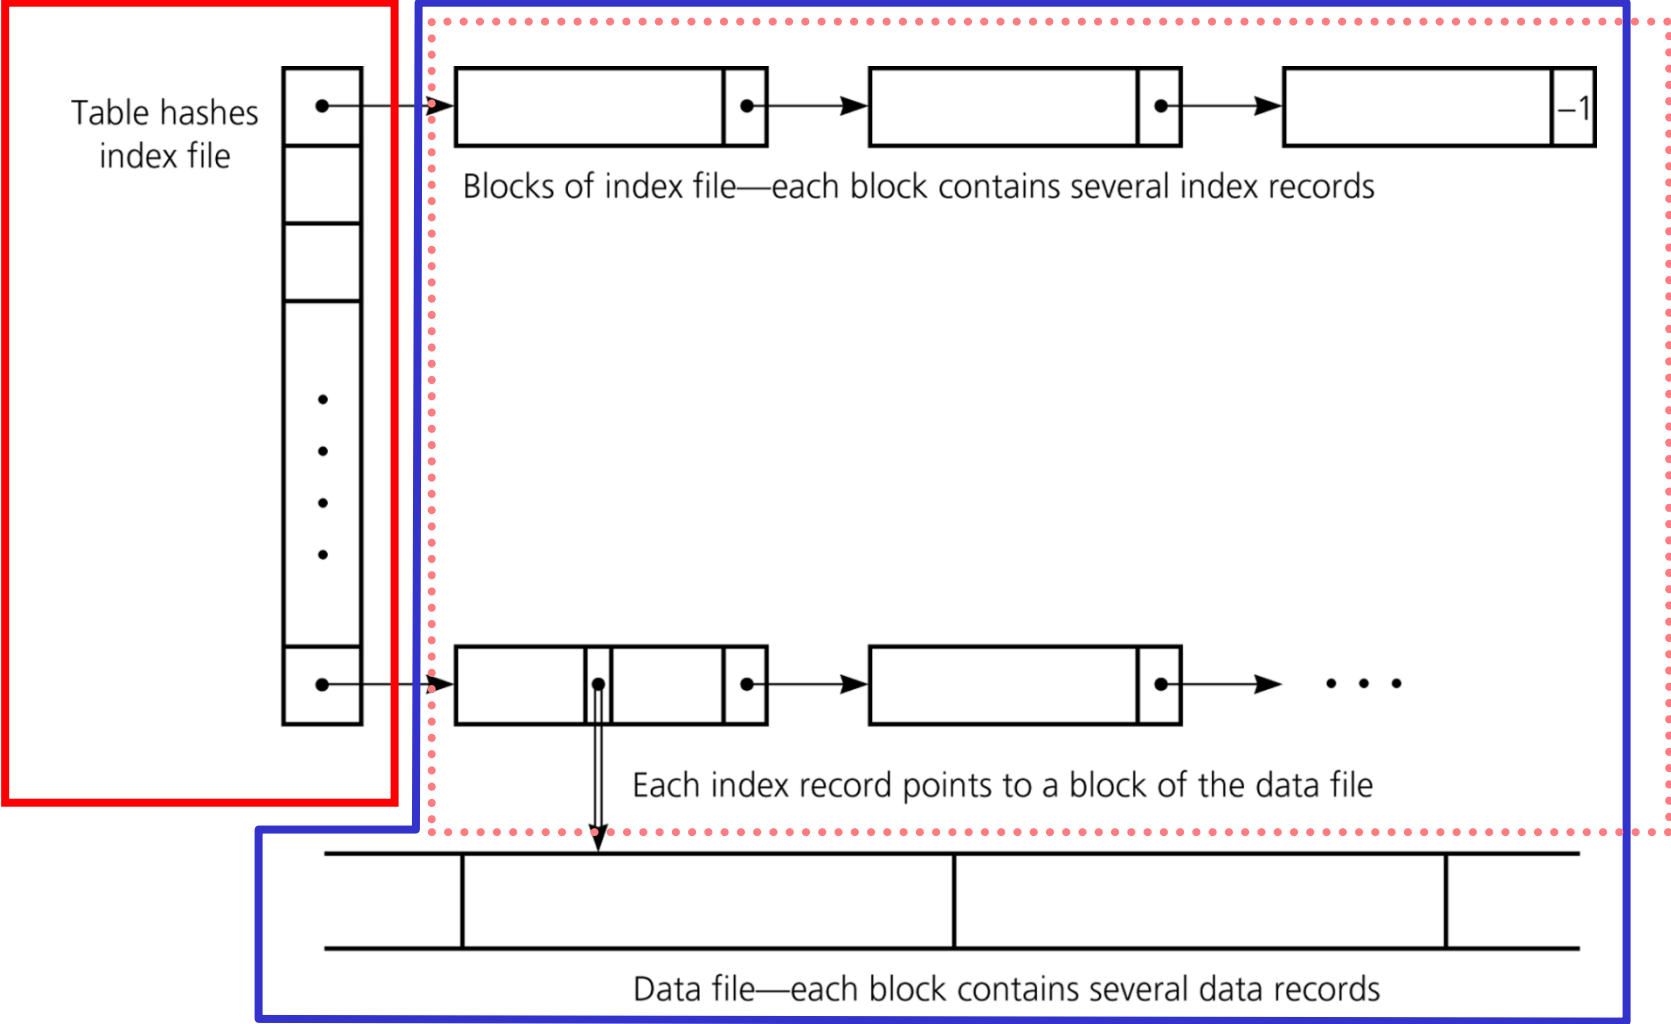
In external table, indices are managed not by references but by block numbers



External Hashing

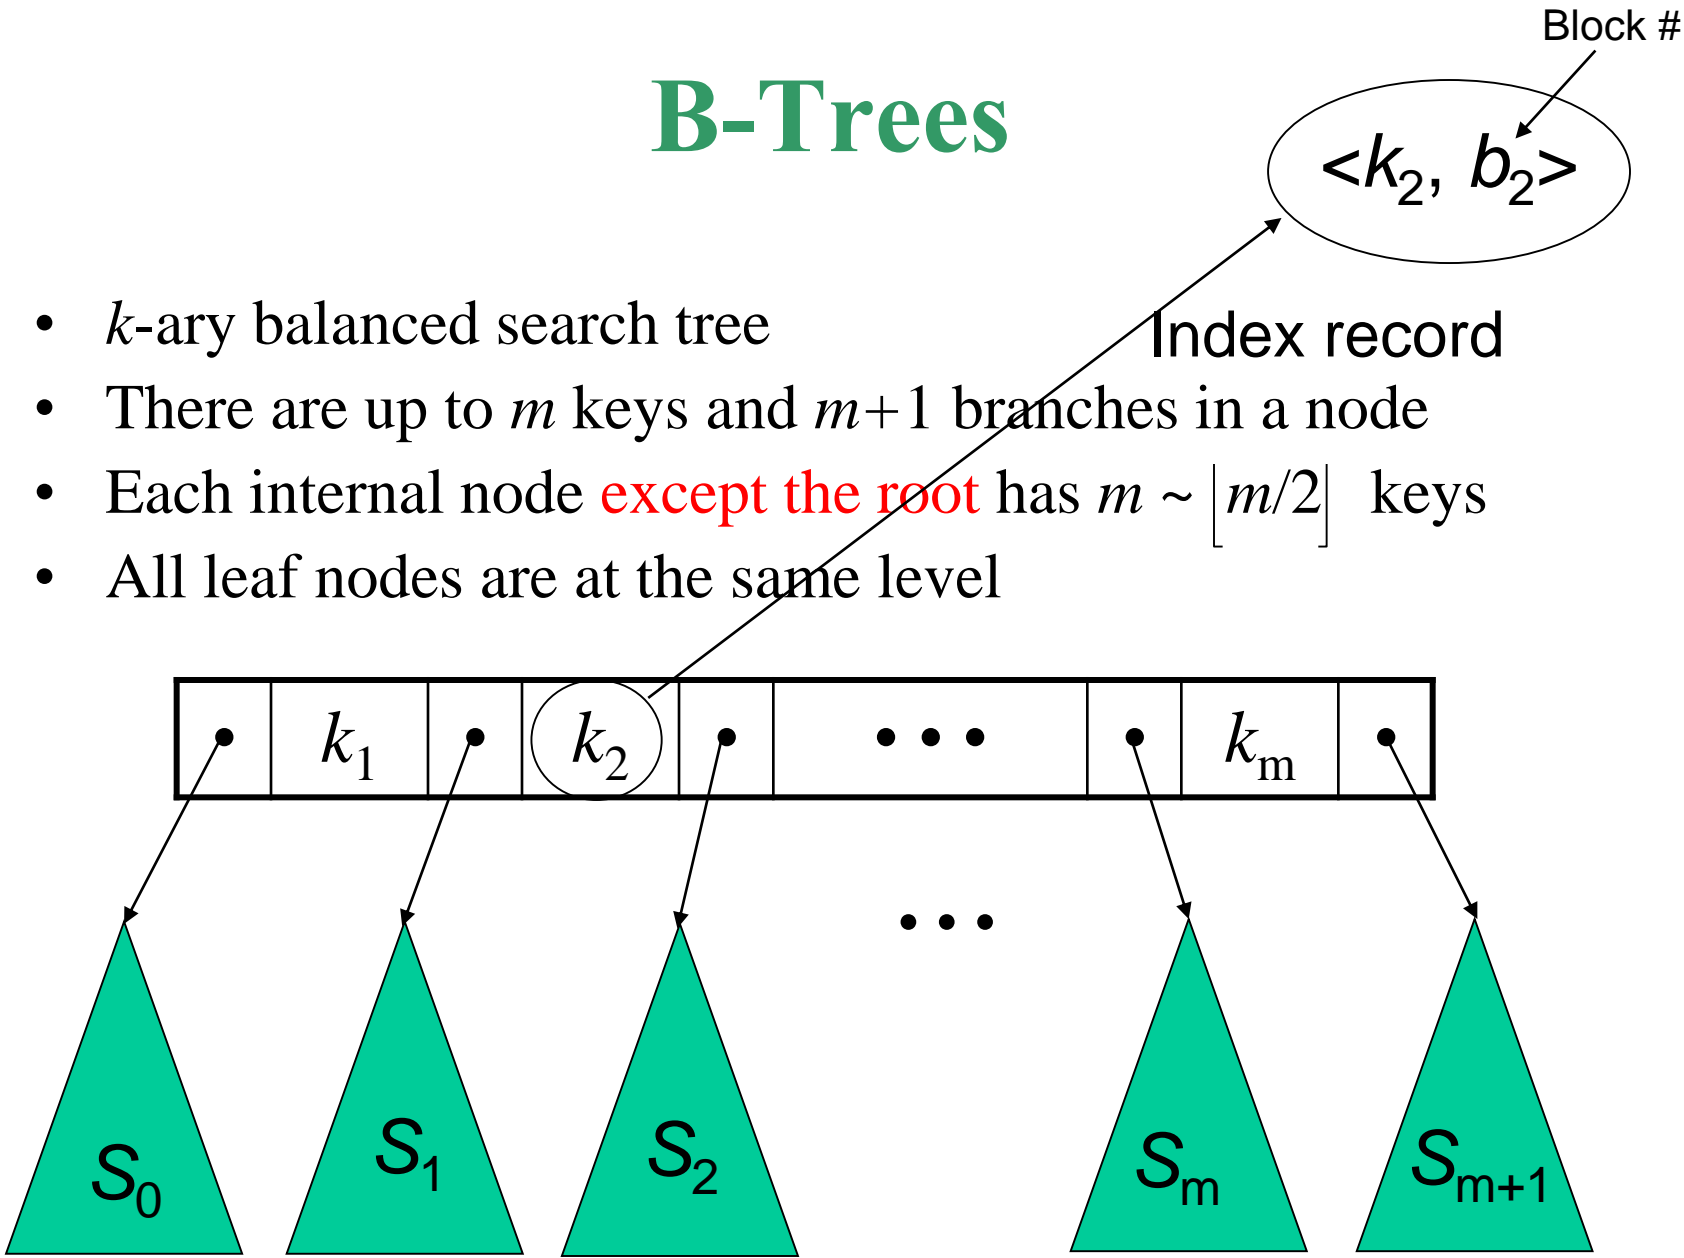
Main memory

Disk



B-Trees

- k -ary balanced search tree
- There are up to m keys and $m+1$ branches in a node
- Each internal node **except the root** has $m \sim \lfloor m/2 \rfloor$ keys
- All leaf nodes are at the same level



Retrieval in a B-tree

BRetrieve(*fIndex*, *fData*, *rootNum*, *key*)

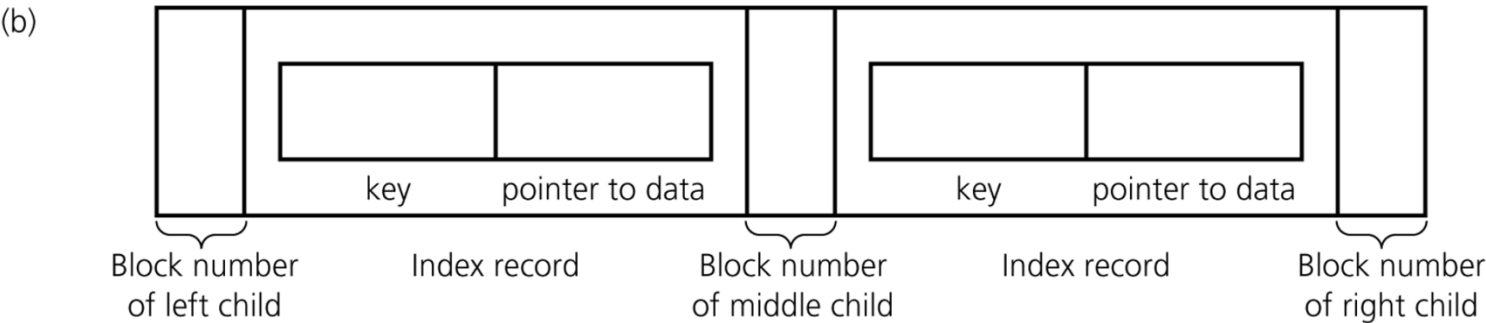
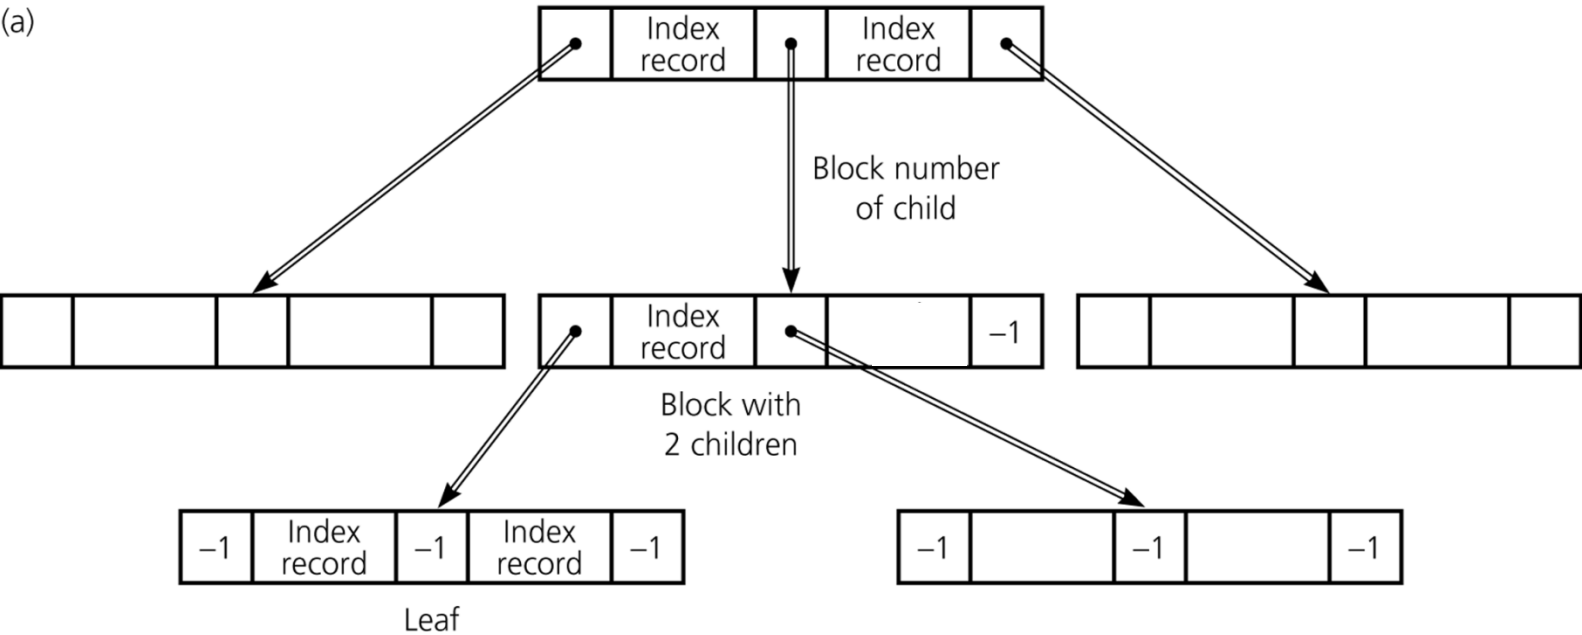
// *fIndex*: B-tree file; *fData*: data file; *key*: search key

// *rootNum*: block # that contains the root of the tree(subtree)

```
{
    if (rootNum == -1) return null;
    buf.readBlock(fIndex, rootNum);
    if (key is one of the  $k_i$ 's in the root) { // get data item
         $b_i$  = data-file block # that index record specifies;
        buf.readBlock(fData,  $b_i$ );
        return the item corresponding to key from buf;
    } else { // branch
         $b_i$  = block # to branch;
        BRetrieve(fIndex, fData,  $b_i$ , key);
    }
}
```

✓ 각 노드가 하나씩의 disk block


Node Structure of a B-tree






Insertion in a B-Tree

- Insert the record w/ key x into the data file
 - Any block p with a vacant slot or a new block
- Insert the index record into the index file
 - Insert the index record $\langle x, p \rangle$ into the B-tree
 - Analogous to the insertion of 2-3 tree

Remind: Insertion in 2-3 Trees


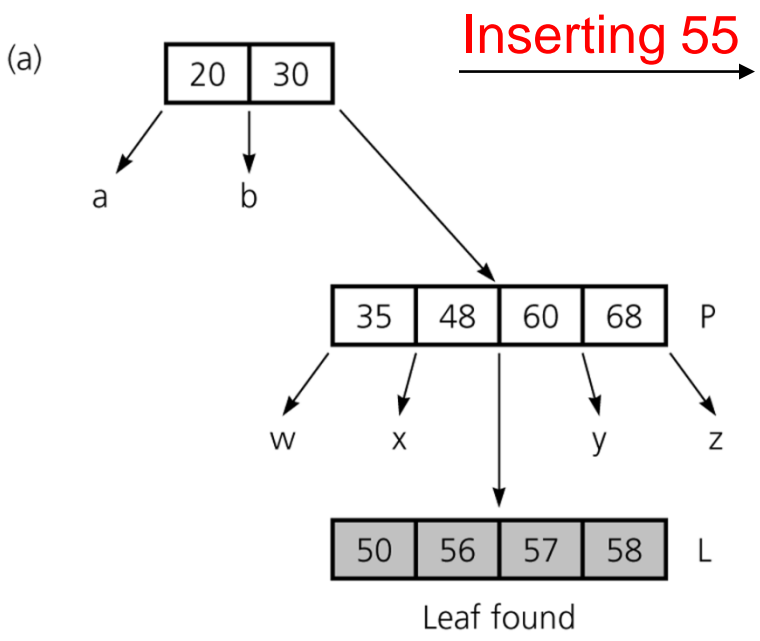




inserting 38



An Example Insertion in a B-Tree


가정: 노드당 최대 4개의 key 허용
(i.e., 2~4개)





- ✓ 이 예는 그냥 split을 했지만
일반적으로는 split보다 redistribution을 먼저 시도한다

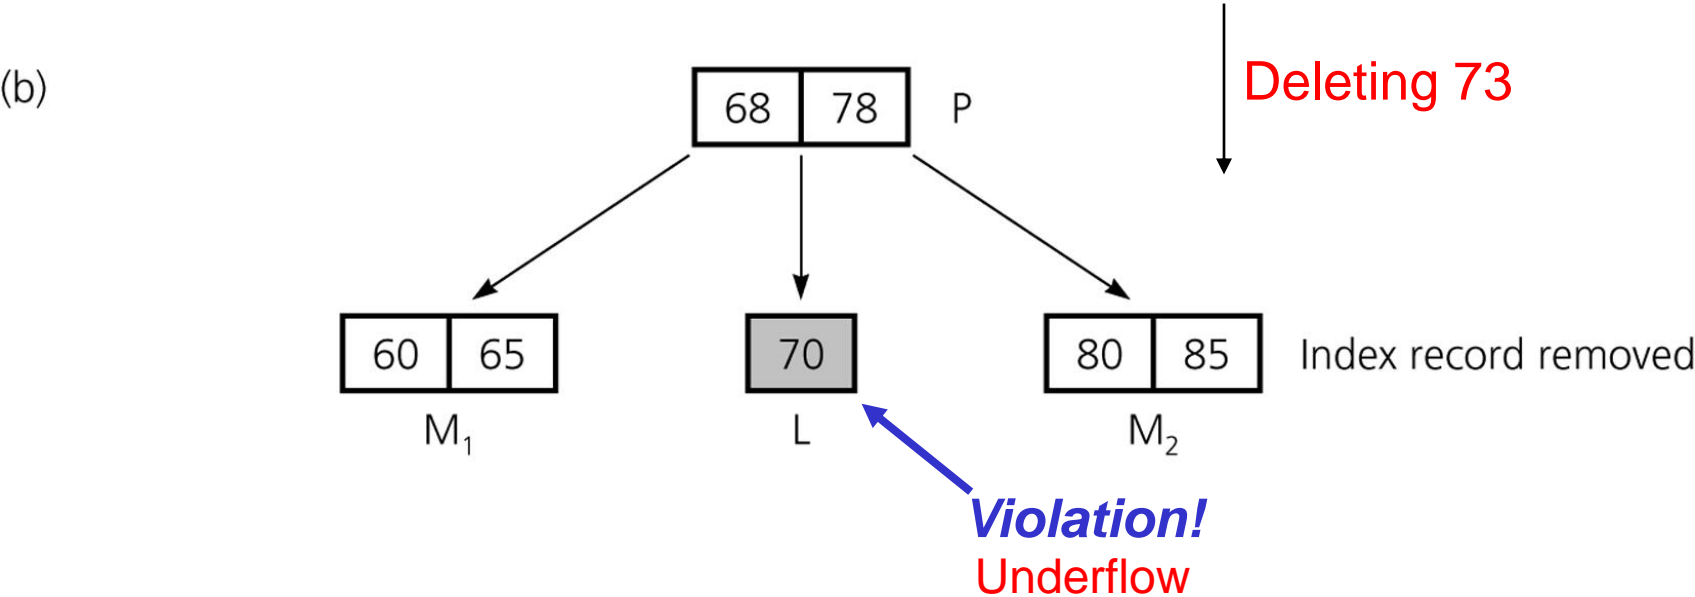

Splitting the Root in a B-Tree



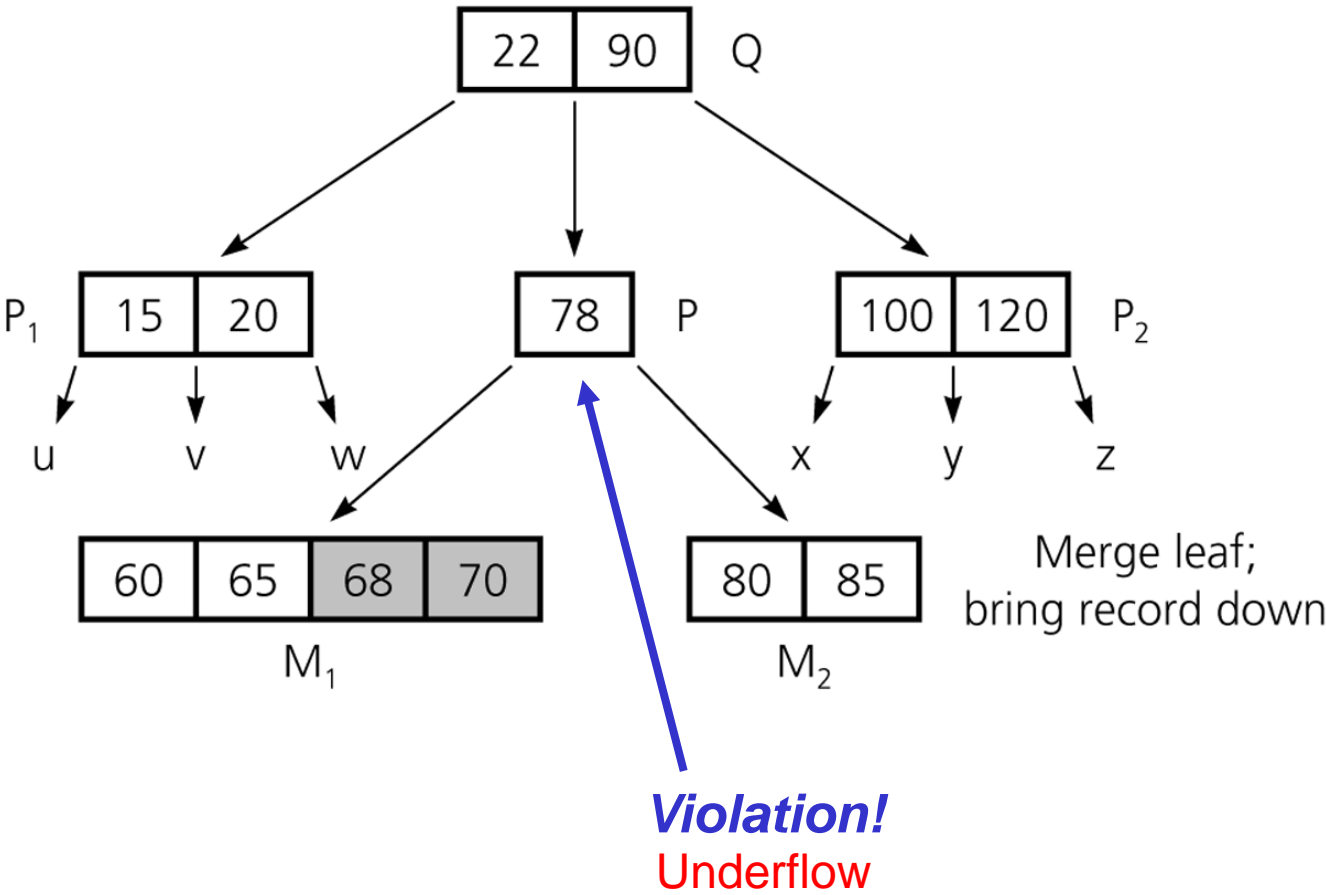
Deletion in a B-Tree

- Find the index record (with the key) and delete in the B-tree
 - Locate the index record
 - If the record is not in a leaf, swap it with its inorder successor
 - Delete the key
 - If the leaf now has fewer than $\lfloor m/2 \rfloor$ keys, merge appropriately
- Delete the corresponding item in the data file

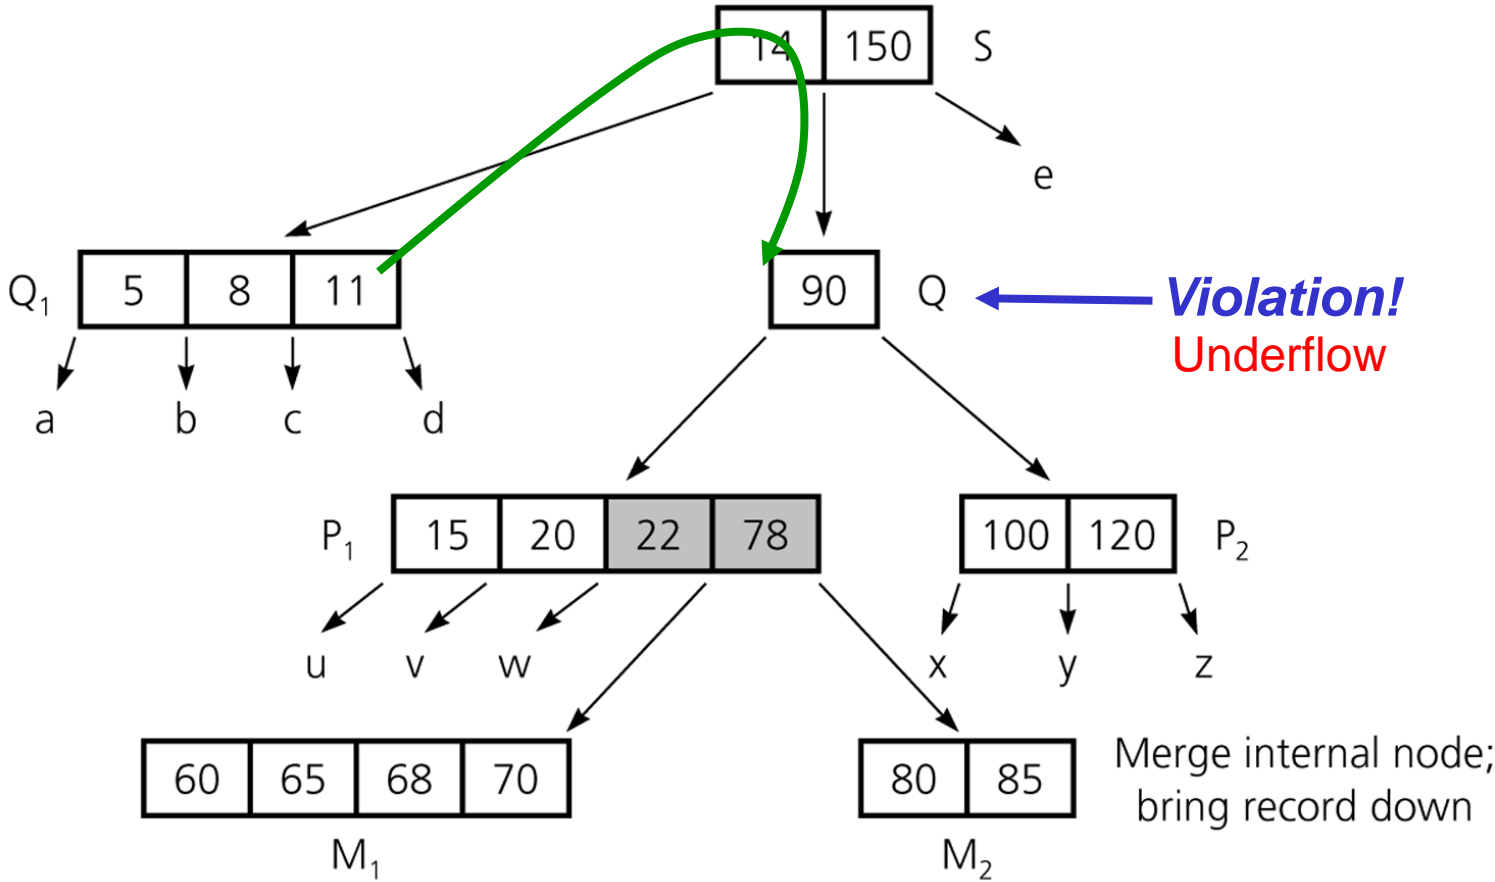
An Example Deletion in a B-Tree




(c)



(d)




(e)



Removing the Root in a B-Tree

Situation: merging w.r.t. the root w/ only one key



Multiple Indexing

